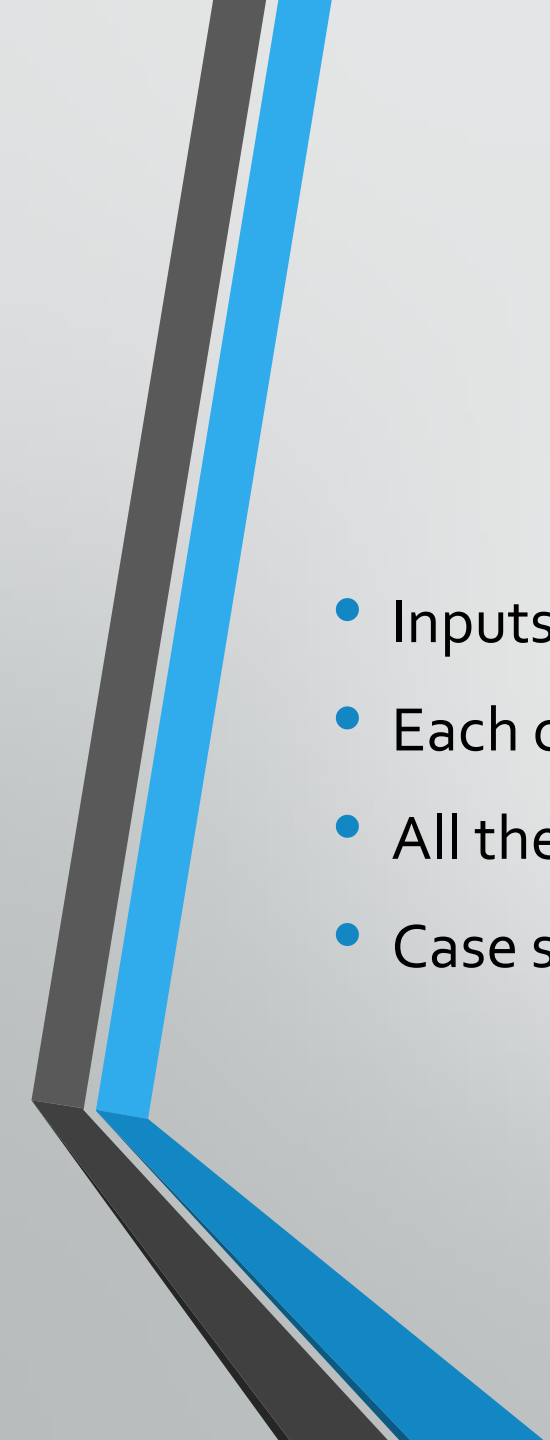




# Lockstep


Study and implementation of a network library for  
distributed deterministic simulations

Raffaele Zippo, Enrico Meloni

- 
- Inputs for the simulation are generated by clients
  - Each client computes the new state of the simulation by applying inputs
  - All the clients need a coherent view of the state
  - Case study: Multiplayer RTS

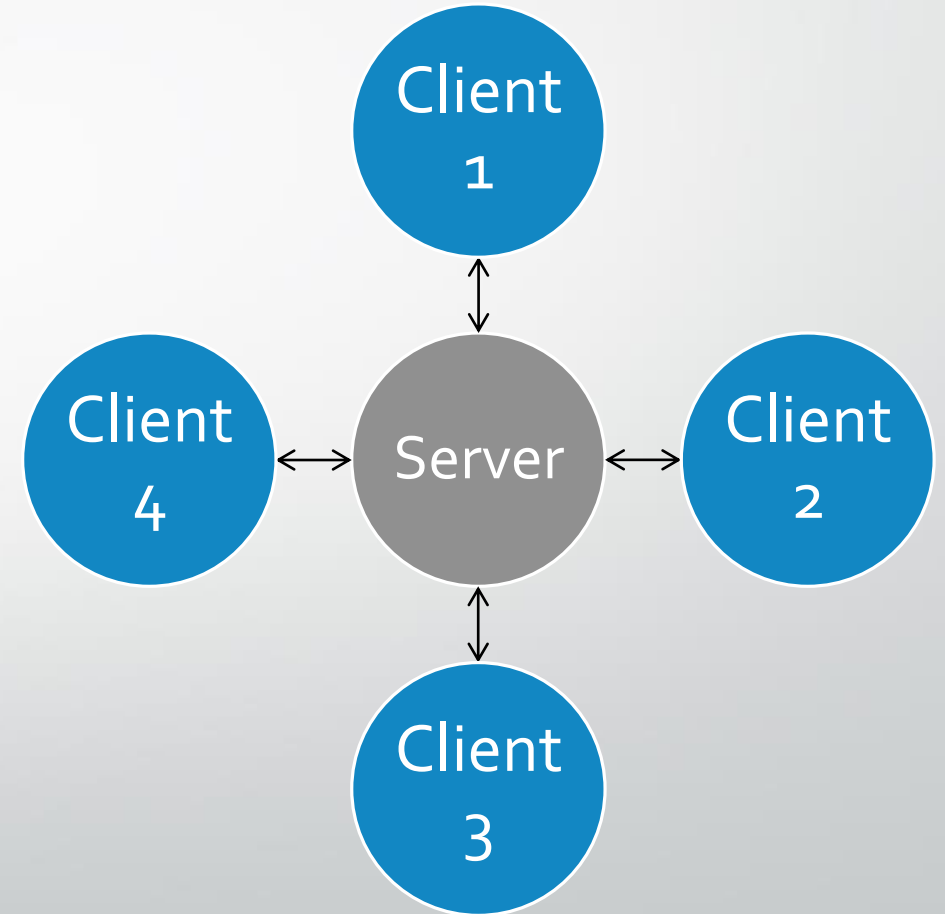
# The FrameQueue

- Time is divided in chunks called *frames*. Each input is assigned to a frame
- All clients execute the same input during the same frame
- By scheduling new inputs for a later frame, we can introduce enough delay to propagate it to all clients and keep executing without interruptions
- The delay introduced should be dimensioned for the application requirements and network characteristics



# Network architecture

- Client-Server architecture
  - Each clients sends its inputs to the server
  - The server forwards to each client others' inputs
- Better performance and scalability
  - Assuming better client-server connection
  - P2P would easily saturate client's upload and increase delays



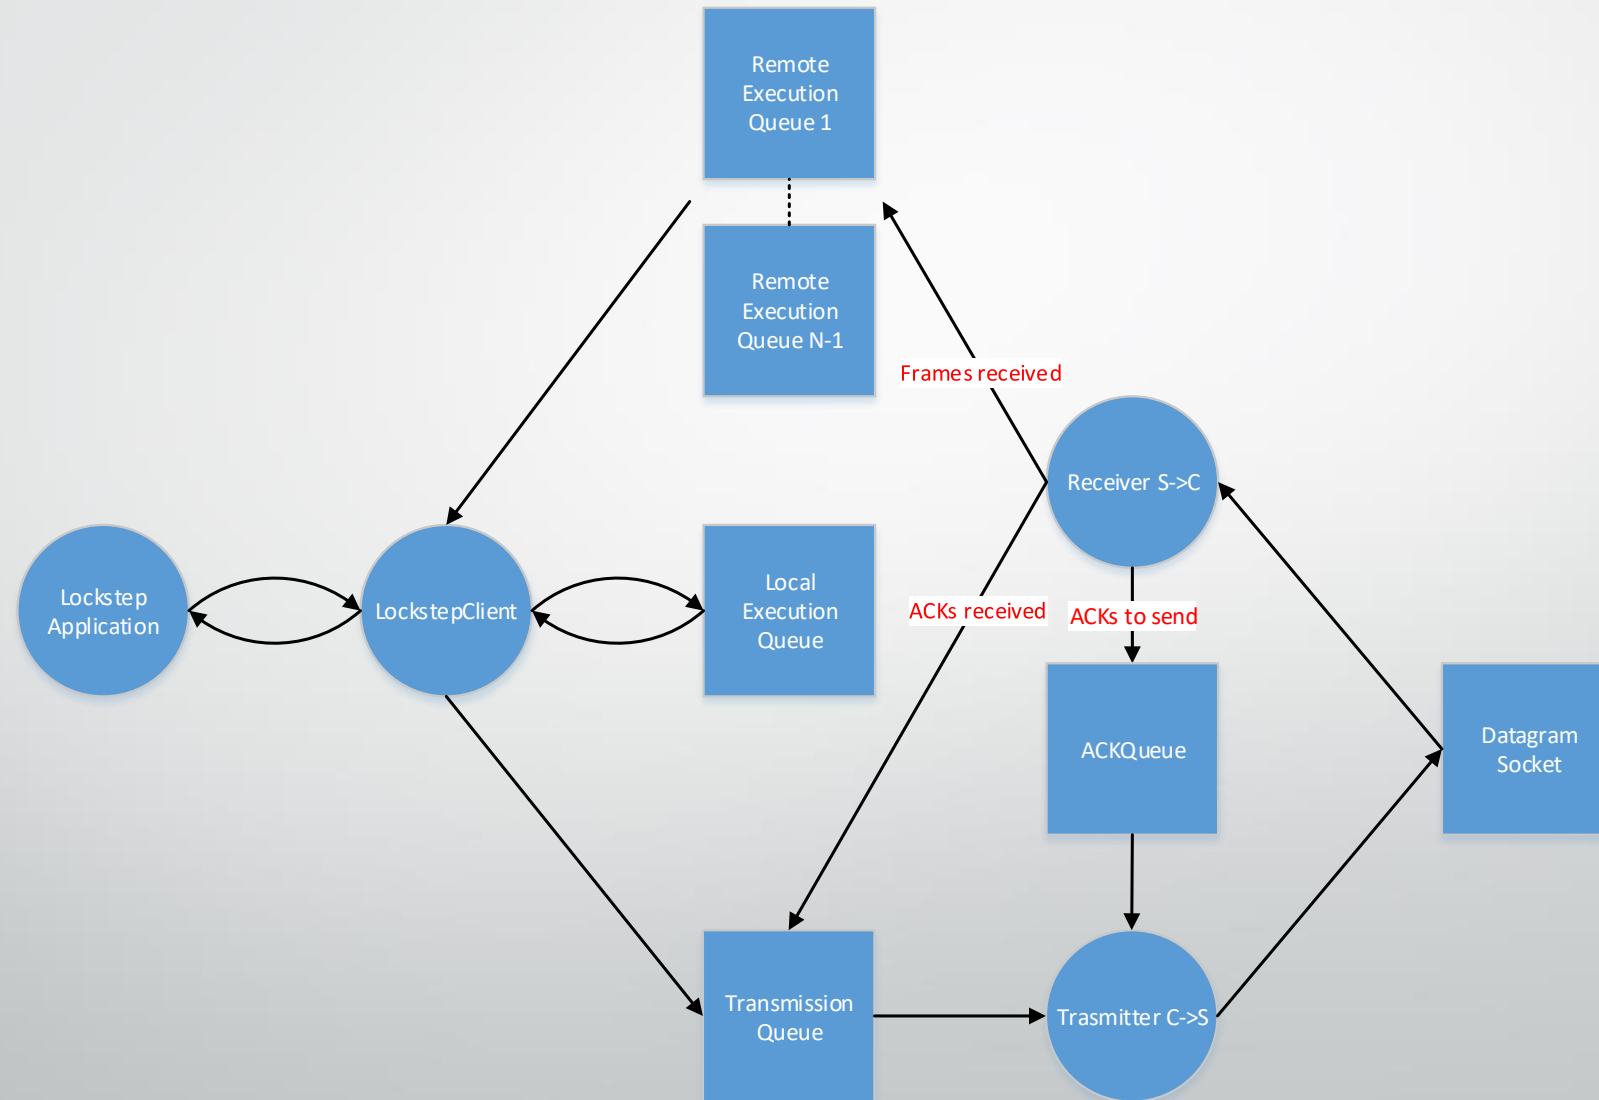
# Network protocol

- Custom protocol based on UDP
  - Reliability
  - Minimum delay
- LockstepCommand: interface implementing Serializable
  - Implemented by the application command
  - Using Java serialization leads to high bandwidth usage
  - Minor solution by adding GZIP compression of outgoing packets

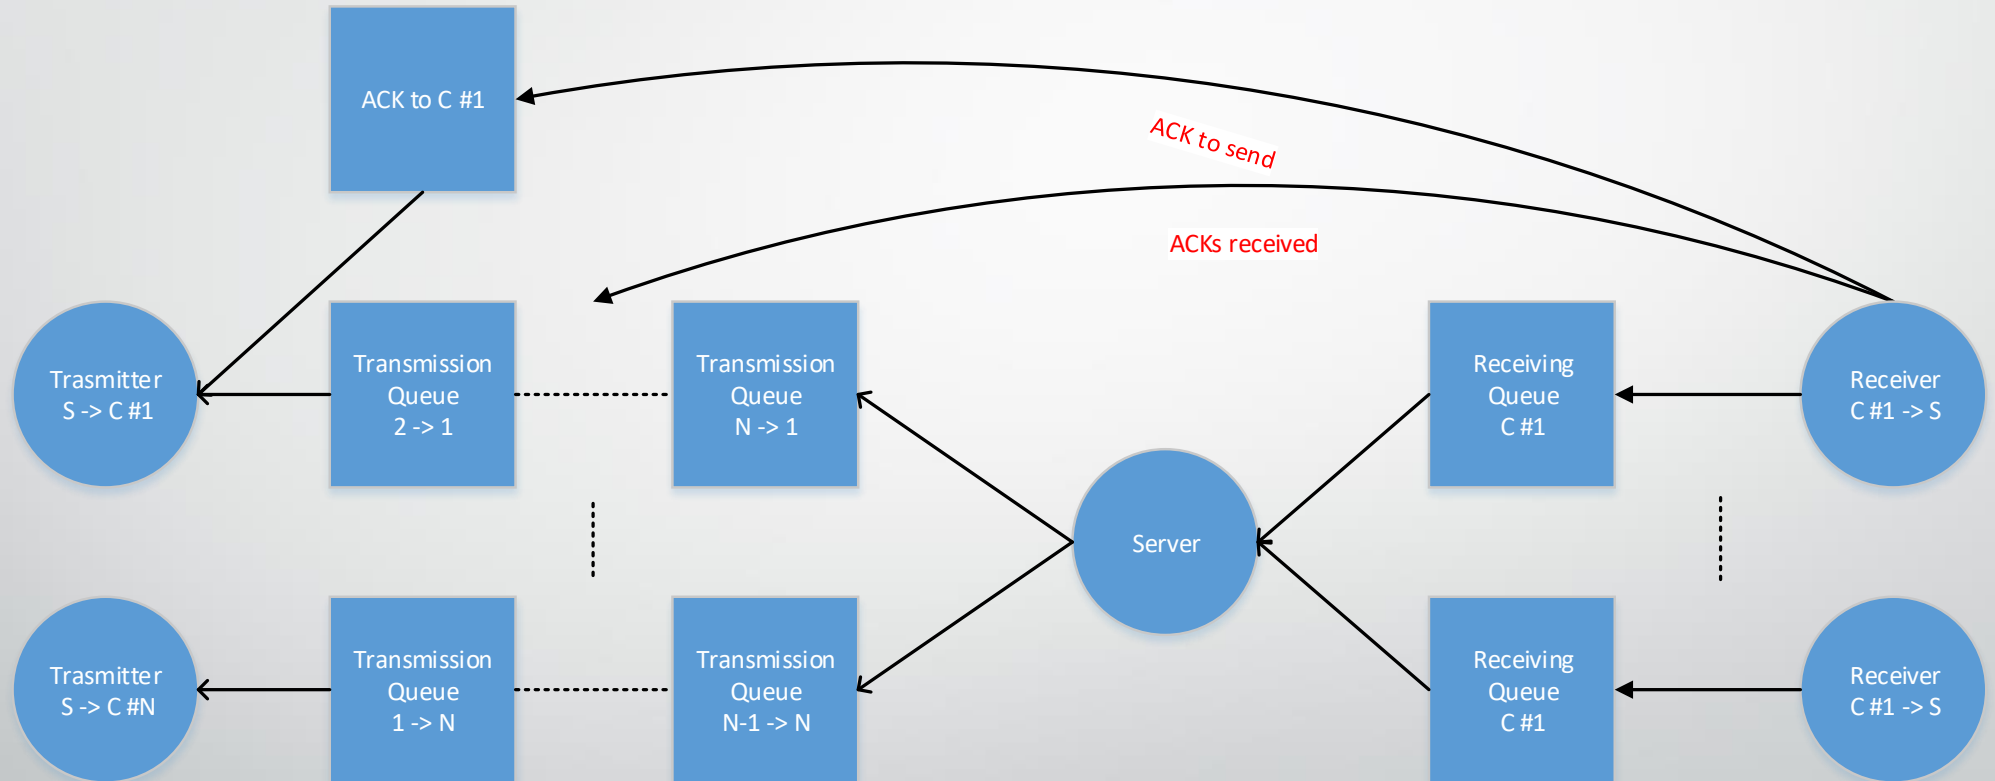


# Software demonstration

# Client Structure



# Server structure





# Client-side Synchronization (1)

- $N$  clients  $C_i$ , each one with  $N$  execution queues  $E_j$
- $E_j$  stores every received input generated by  $C_j$
- $C_i$  and  $E_j$  keep track of the next frame number, let's say  $k$
- $C_i$  needs command  $k$  from every  $E_j$  to execute frame  $k$
- If some command is missing in its queue, client stops and wait

# Client-side Synchronization (2)


- Synchronization is achieved with a Semaphore
- $E_j$  releases a permit when it has the next command  $k$
- $C_i$  tries to acquire  $N$  permits from semaphore
  - Failure to acquire: signal the application to suspend simulation, wait for all permits
  - Permits available: continue execution, resume simulation signal if suspended

# Server-side Synchronization

- $N$  execution queues  $E_j$  for storing input received by  $C_j$
- Semaphore is used
- $E_j$  releases a permit as soon as it receives a command
- Server waits on semaphore
- When input is available, server forwards it to transmitters
- Clients wait for in-order command, server waits for any command

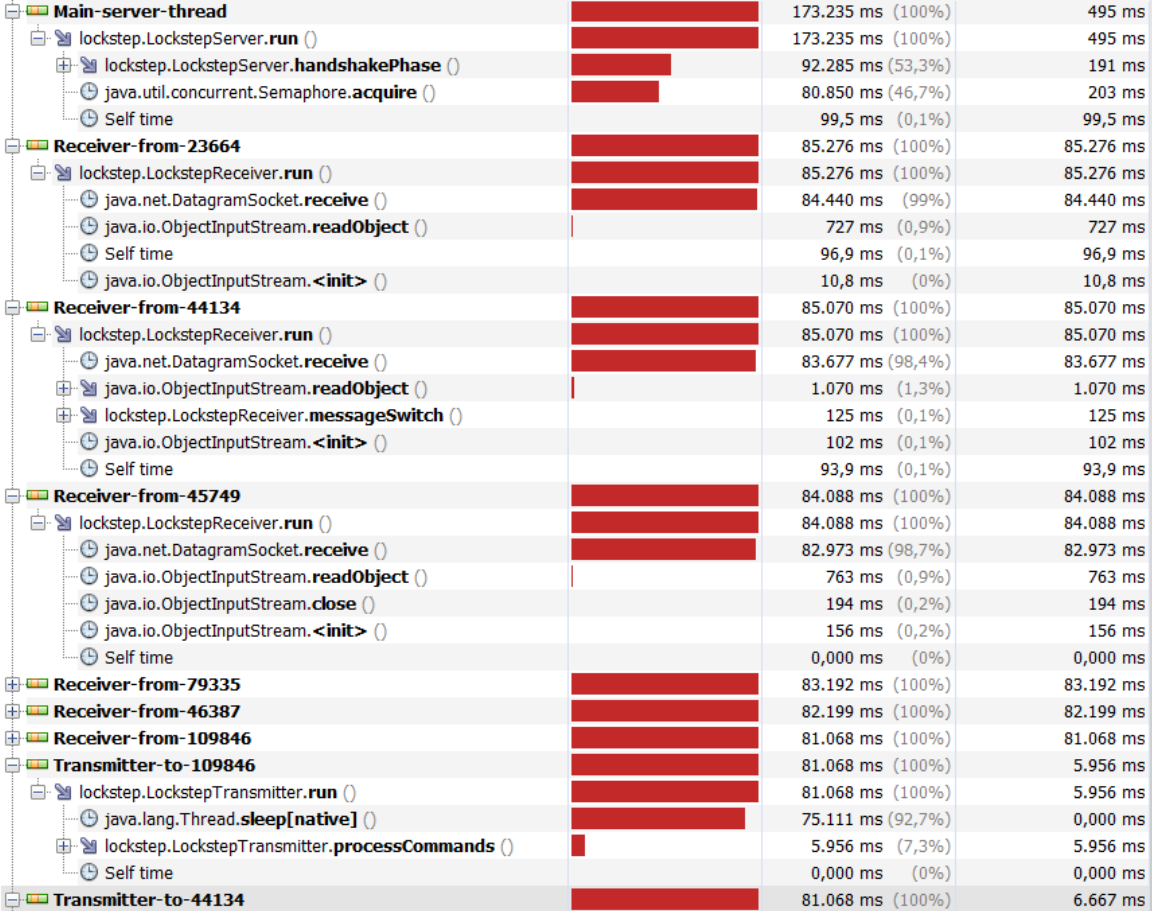
# Scalability

- The protocol has low bandwidth needs for clients
  - Receives  $O(N)$  messages
  - Sends  $O(1)$  messages
- Server instead has high bandwidth needs
  - Receives  $O(N)$  messages
  - Sends  $O(N^2)$  messages



# Future developments

- Receivers and transmitters are mostly idle
- Use idle time to improve serialization
  - Key objective: reducing bandwidth
- Optional P2P architecture
  - For high-bandwidth clients (LAN)
- Adaptive input-to-execution delay
- Security
  - Authentication
  - Integrity
  - State validation (anti-cheat)



Main-server-thread		173.235 ms (100%)	495 ms
lockstep.LockstepServer.run ()		173.235 ms (100%)	495 ms
lockstep.LockstepServer.handshakePhase ()		92.285 ms (53,3%)	191 ms
java.util.concurrent.Semaphore.acquire ()		80.850 ms (46,7%)	203 ms
Self time		99,5 ms (0,1%)	99,5 ms
Receiver-from-23664		85.276 ms (100%)	85.276 ms
lockstep.LockstepReceiver.run ()		85.276 ms (100%)	85.276 ms
java.net.DatagramSocket.receive ()		84.440 ms (99%)	84.440 ms
java.io.ObjectInputStream.readObject ()		727 ms (0,9%)	727 ms
Self time		96,9 ms (0,1%)	96,9 ms
java.io.ObjectInputStream.<init> ()		10,8 ms (0%)	10,8 ms
Receiver-from-44134		85.070 ms (100%)	85.070 ms
lockstep.LockstepReceiver.run ()		85.070 ms (100%)	85.070 ms
java.net.DatagramSocket.receive ()		83.677 ms (98,4%)	83.677 ms
java.io.ObjectInputStream.readObject ()		1.070 ms (1,3%)	1.070 ms
lockstep.LockstepReceiver.messageSwitch ()		125 ms (0,1%)	125 ms
java.io.ObjectInputStream.<init> ()		102 ms (0,1%)	102 ms
Self time		93,9 ms (0,1%)	93,9 ms
Receiver-from-45749		84.088 ms (100%)	84.088 ms
lockstep.LockstepReceiver.run ()		84.088 ms (100%)	84.088 ms
java.net.DatagramSocket.receive ()		82.973 ms (98,7%)	82.973 ms
java.io.ObjectInputStream.readObject ()		763 ms (0,9%)	763 ms
java.io.ObjectInputStream.close ()		194 ms (0,2%)	194 ms
java.io.ObjectInputStream.<init> ()		156 ms (0,2%)	156 ms
Self time		0,000 ms (0%)	0,000 ms
Receiver-from-79335		83.192 ms (100%)	83.192 ms
Receiver-from-46387		82.199 ms (100%)	82.199 ms
Receiver-from-109846		81.068 ms (100%)	81.068 ms
Transmitter-to-109846		81.068 ms (100%)	5.956 ms
lockstep.LockstepTransmitter.run ()		81.068 ms (100%)	5.956 ms
java.lang.Thread.sleep[native] ()		75.111 ms (92,7%)	0,000 ms
lockstep.LockstepTransmitter.processCommands ()		5.956 ms (7,3%)	5.956 ms
Self time		0,000 ms (0%)	0,000 ms
Transmitter-to-44134		81.068 ms (100%)	6.667 ms